

**comparative execution time as well as analysis of the performance of different kernels (how many times faster, why is it faster, overheads in calling the kernel, etc.)**

The programming language and number of threads that we use affects the run time of the kernel in our stencil vector program.

	2 <sup>20</sup>	2 <sup>24</sup>	2 <sup>28</sup>
C	125000 microseconds	2002000 microseconds	31359000 microseconds
CUDA	85.800 microseconds	765.78 microseconds	13448 microseconds
x86-64	83000 microseconds	1327000 microseconds	20797000 microseconds
SIMD	14000 microseconds	268000 microseconds	3912000 microseconds

As we can see on the table above, CUDA performed the fastest but this is due to having more workers or threads to run the kernel at the same time. In our CUDA program we ran with 1024 threads and blocks for all input sizes which certainly gave it an advantage over C, x86-62 and SIMD. For the input size 2<sup>20</sup> CUDA kernel ran 1457 times faster than C, 967 times faster than x86-64, and 163 times faster than SIMD. For the input size 2<sup>24</sup> CUDA kernel ran 2614 times faster than C, 1732 times faster than x86-64, and 350 times faster than SIMD. For the input size 2<sup>28</sup> CUDA kernel ran 2331 times faster than C, 1546 times faster than x86-64, and 290 times faster than SIMD. However, CUDA still has some overheads due to the transfer of host to device and vice-versa.

The second fastest would be SIMD. This is due to the fact that the language is a low level programming language which is close to machine language as well as it has a certain degree of parallelism. In SIMD we can use `shr <array size>, <value>` which shifts the bits to the right. This allows the array to be divided into 2 raised to <value>. In our case the value we used is 3, thus our array was divided into 8 and processed at the same time. For input size 2<sup>20</sup> SIMD performed almost 6 times faster than x86-64 and almost 9 times faster than C. For input size 2<sup>24</sup> SIMD performed almost 5 times faster than x86-64, and 74 times faster than C.

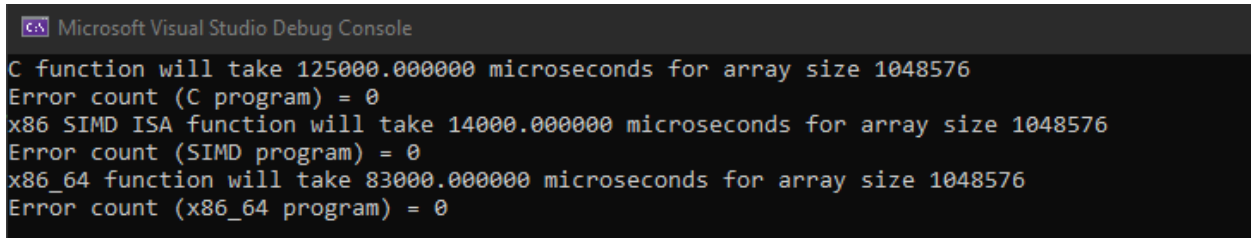
Following SIMD in terms of runtime would be x86-64. This is due to x86-62 being a low level programming language which makes it close to machine language which means there's no need for translation or interpretation. However in x86-64 we processed the program sequentially instead of parallel thus it is slower than SIMD. For input size  $2^{20}$ ,  $2^{24}$  and  $2^{28}$  x86-62 performed approximately 1.5 times faster than C.

Lastly C had the slowest run time of the kernel compared to all the programming languages discussed above. This is due to C being a high level programming language which means it still needs translation or interpretation for it to be understood by the machine which makes it slower. Additionally, we also performed the kernel sequentially wherein we did not have any threads to work on the kernel in parallel.

In conclusion, the programming language and parallelism affects the run time of programs. Parallelism improves the efficiency of a program by allowing threads or workers to work on the kernel. The more threads we have the faster the kernel would be processed. However we should not mindlessly increase thread counts as it eats more resources, possible race conditions, and we should be mindful of the block size to optimize the number of threads. Moreover low level programming languages run faster than high level programming languages due to the run time consumed for the translation of high level programming languages to commands our machine can understand.

### **Screenshot of the program output with correctness check C, x86-64, SIMD YMM register**

For input size  $2^{20}$ :



```
Microsoft Visual Studio Debug Console
C function will take 125000.000000 microseconds for array size 1048576
Error count (C program) = 0
x86 SIMD ISA function will take 14000.000000 microseconds for array size 1048576
Error count (SIMD program) = 0
x86_64 function will take 83000.000000 microseconds for array size 1048576
Error count (x86_64 program) = 0
```

For input size  $2^{24}$ :

#### Microsoft Visual Studio Debug Console

```
C function will take 2002000.000000 microseconds for array size 16777216
Error count (C program) = 0
x86 SIMD ISA function will take 268000.000000 microseconds for array size 16777216
Error count (SIMD program) = 0
x86_64 function will take 1327000.000000 microseconds for array size 16777216
Error count (x86_64 program) = 0
```

For input size  $2^{28}$ :

#### Microsoft Visual Studio Debug Console

```
C function will take 31359000.000000 microseconds for array size 268435456
Error count (C program) = 0
x86 SIMD ISA function will take 3912000.000000 microseconds for array size 268435456
Error count (SIMD program) = 0
x86_64 function will take 20797000.000000 microseconds for array size 268435456
Error count (x86_64 program) = 0
```

### Screenshot of the program output including correctness check (CUDA, optional)

For input size  $2^{20}$ :

```
1 %shell
2 nvprof ./c_stencilVector

==1839== NVPROF is profiling process 1839, command: ./c_stencilVector
numBlocks = 1024 numThreads = 1024
Error count(CUDA program): 0
==1839== Profiling application: ./c_stencilVector
==1839== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 100.00%  2.5740ms      30  85.800us  85.215us  86.559us  stencilVector(int, int*, int*)
API calls: 97.64%  245.85ms      2  122.92ms  36.132us  245.81ms  cudaMallocManaged
1.02%  2.5581ms      1  2.5581ms  2.5581ms  2.5581ms  cudaDeviceSynchronize
0.98%  2.4597ms      4  614.92us  163.78us  1.1794ms  cudaMemPrefetchAsync
0.24%  603.58us      2  301.79us  283.86us  319.72us  cudaFree
0.06%  156.44us      30  5.2140us  3.3750us  41.309us  cudaLaunchKernel
0.05%  117.72us      101  1.1650us  140ns  49.367us  cuDeviceGetAttribute
0.01%  25.187us      1  25.187us  25.187us  25.187us  cuDeviceGetName
0.01%  14.026us      2  7.0130us  1.9290us  12.097us  cudaMemAdvise
0.00%  6.4470us      1  6.4470us  6.4470us  6.4470us  cuDeviceGetPCIBusId
0.00%  1.9170us      1  1.9170us  1.9170us  1.9170us  cudaGetDevice
0.00%  1.7690us      3  589ns  189ns  1.2760us  cuDeviceGetCount
0.00%  945ns      2  472ns  207ns  738ns  cuDeviceGet
0.00%  544ns      1  544ns  544ns  544ns  cuModuleGetLoadingMode
0.00%  400ns      1  400ns  400ns  400ns  cuDeviceTotalMem
0.00%  227ns      1  227ns  227ns  227ns  cuDeviceGetUuid

==1839== Unified Memory profiling result:
Device "Tesla T4 (0)"
   Count   Avg Size   Min Size   Max Size   Total Size   Total Time   Name
2  2.0000MB  2.0000MB  2.0000MB  4.000000MB  355.4860us  Host To Device
2  2.0000MB  2.0000MB  2.0000MB  4.000000MB  321.9800us  Device To Host
```

For input size  $2^{24}$ :

```
1 %shell
2 nvprof ./c_stencilVector

==3404== NVPROF is profiling process 3404, command: ./c_stencilVector
numBlocks = 1024 numThreads = 1024
Error count(CUDA program): 0
==3404== Profiling application: ./c_stencilVector
==3404== Profiling result:
   Type  Time(%)    Time       Calls      Avg        Min        Max   Name
GPU activities: 100.00%  22.973ms    30      765.78us    762.78us    769.18us stencilVector(int, int*, int*)
API calls:      78.26%  244.43ms    2      122.21ms    50.465us    244.38ms cudaMallocManaged
              11.06%  34.553ms    4      8.6382ms    404.42us    18.224ms cudaMemPrefetchAsync
              7.32%   22.855ms    1      22.855ms    22.855ms    22.855ms cudaDeviceSynchronize
              3.24%   10.118ms    2      5.0592ms    4.6176ms    5.5008ms cudaFree
              0.06%   199.15us    30      6.6380us    3.3200us    54.682us cudaLaunchKernel
              0.04%   111.41us   101      1.1030us    133ns      46.131us cuDeviceGetAttribute
              0.01%   25.423us    1      25.423us    25.423us    25.423us cuDeviceGetName
              0.01%   15.854us    2      7.9270us    1.8370us    14.017us cudaMemAdvise
              0.00%   5.6010us    1      5.6910us    5.6910us    5.6910us cuDeviceGetPCIBusId
              0.00%   2.5940us    1      2.5940us    2.5940us    2.5940us cudaGetDevice
              0.00%   1.6270us    3         542ns     190ns     1.1140us cuDeviceGetCount
              0.00%    991ns     2         495ns     297ns      694ns cuDeviceGet
              0.00%    487ns     1         487ns     487ns     487ns cuModuleGetLoadingMode
              0.00%    409ns     1         409ns     409ns     409ns cuDeviceTotalMem
              0.00%    267ns     1         267ns     267ns     267ns cuDeviceGetUuid

==3404== Unified Memory profiling result:
Device "Tesla T4 (0)"
  Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    32    2.0000MB  2.0000MB  2.0000MB  64.00000MB  5.629438ms  Host To Device
    32    2.0000MB  2.0000MB  2.0000MB  64.00000MB  5.167804ms  Device To Host
```

For input size 2^28:

```
1 %shell
2 nvprof ./c_stencilVector

==6172== NVPROF is profiling process 6172, command: ./c_stencilVector
numBlocks = 1024 numThreads = 1024
Error count(CUDA program): 0
==6172== Profiling application: ./c_stencilVector
==6172== Profiling result:
   Type  Time(%)    Time       Calls      Avg        Min        Max   Name
GPU activities: 100.00%  403.45ms    30     13.448ms    13.184ms    14.042ms stencilVector(int, int*, int*)
API calls:      37.83%  553.99ms    4     138.50ms    7.8323ms    293.15ms cudaMemPrefetchAsync
              27.54%  403.34ms    1     403.34ms    403.34ms    403.34ms cudaDeviceSynchronize
              17.46%  255.62ms    2     127.81ms    124.96ms    130.66ms cudaFree
              17.15%  251.11ms    2     125.55ms    48.883us    251.06ms cudaMallocManaged
              0.01%   201.30us    30      6.7090us    3.2000us    59.132us cudaLaunchKernel
              0.01%   112.23us   101      1.1110us    130ns      46.680us cuDeviceGetAttribute
              0.00%   25.267us    1      25.267us    25.267us    25.267us cuDeviceGetName
              0.00%   13.715us    2      6.8570us    2.0060us    11.709us cudaMemAdvise
              0.00%   5.9550us    1      5.9550us    5.9550us    5.9550us cuDeviceGetPCIBusId
              0.00%   1.8410us    1      1.8410us    1.8410us    1.8410us cudaGetDevice
              0.00%   1.6070us    3         535ns     198ns     1.1110us cuDeviceGetCount
              0.00%   1.1040us    2         552ns     298ns      806ns cuDeviceGet
              0.00%    496ns     1         496ns     496ns     496ns cuModuleGetLoadingMode
              0.00%    423ns     1         423ns     423ns     423ns cuDeviceTotalMem
              0.00%    236ns     1         236ns     236ns     236ns cuDeviceGetUuid

==6172== Unified Memory profiling result:
Device "Tesla T4 (0)"
  Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    512    2.0000MB  2.0000MB  2.0000MB  1.000000GB  89.77512ms  Host To Device
    512    2.0000MB  2.0000MB  2.0000MB  1.000000GB  82.17858ms  Device To Host
```